

Visualizing United States Census Data

Eli Turkel
Nadia Antony

5/21/2019

Agenda

- History of the U.S. Census
- Getting started with Tidycensus
- Concepts - Tidy Data
- Step by step analysis

Follow along to create some cool maps!

US Census Timeline

- 1790: 6 question survey / U.S. pop just under 4 million
- 1810: Census adds data collection on U.S. manufacturers
- 1849: Congress establishes Census Board
- 1890: Herman Hollerith's tabulation machine used in data collection
- 1900: Census becomes a permanent agency and inter-decennial data collection begins
- 1920: For first time, majority of U.S. pop lives in urban areas / U.S. pop over 106 million
- 1940: First use of statistical sampling
- 1960: DIME (Dual Independent Map Encoding)
- 1970: All data products available on magnetic computer tape
- 1980: Mailout / Mailin Surveys
- 2000: U.S. Census asks ten questions
- 2010: American Community Survey 5-year estimates released

Tidycensus Overview

- Developed by Kyle Walker
- [Check out the package on Github!](#)
- Tidycensus is an API wrapper that allows R users to ingest decennial surveys and the American Community Survey (ACS) estimates
- Census data back to the 1990 decennial survey is available
- The default is the 5-year ACS estimate from 2013-2017

Getting an API key

- https://api.census.gov/data/key_signup.html
- You'll then receive an e-mail from the census with your API key

```
api_key <- "xxxxxxxxx"  
census_api_key(api_key, install = TRUE)  
Sys.getenv("CENSUS_API_KEY")
```

Variable Search

```
library(tidycensus)
v17 <- load_variables(2017, "acs5", cache = TRUE)
v17
```

```
## # A tibble: 25,070 x 3
##   name      label      concept
##   <chr>    <chr>    <chr>
## 1 B00001_001 Estimate!!Total UNWEIGHTED SAMPLE COUNT OF THE POPULATION
## 2 B00002_001 Estimate!!Total UNWEIGHTED SAMPLE HOUSING UNITS
## 3 B01001_001 Estimate!!Total SEX BY AGE
## 4 B01001_002 Estimate!!Total!!Male SEX BY AGE
## 5 B01001_003 Estimate!!Total!!Male!!Under 5 years SEX BY AGE
## 6 B01001_004 Estimate!!Total!!Male!!5 to 9 years SEX BY AGE
## 7 B01001_005 Estimate!!Total!!Male!!10 to 14 years SEX BY AGE
## 8 B01001_006 Estimate!!Total!!Male!!15 to 17 years SEX BY AGE
## 9 B01001_007 Estimate!!Total!!Male!!18 and 19 years SEX BY AGE
## 10 B01001_008 Estimate!!Total!!Male!!20 years SEX BY AGE
## # ... with 25,060 more rows
```


Load Var Function

```
vt <- get_acs(geography = "county",  
              variables = c(medincome = "B19013_001"),  
              state = "VT")
```

What is Tidy Data?

Three interrelated rules:

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.



<https://r4ds.had.co.nz/tidy-data.html#fig:tidy-structure>

Why should I care?

R is a vectorized language, meaning you can do operations like:

```
v1 <- c(1, 2, 3, 4)
v2 <- c(5, 6, 7, 8)

v2 - v1
```

```
## [1] 4 4 4 4
```

instead of writing a *for loop* to subtract the individual elements

The packages inside the tidyverse, e.g `dplyr`, let you do data cleaning and manipulation operations easily when data is in tidy format.

Using these packages can help - write faster and more 'readable' code

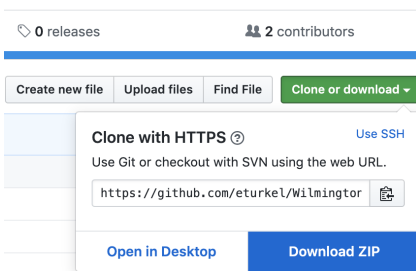
Brief concepts - commands you'll see

- `%>%` the “pipe” operator for chaining
- **`filter`** to subset a dataframe
- **`group_by`** and then **`summarise`**
- **`facet`** to create *small multiples* plots
- **`left_join`** to join datasets
- and `<-` is the assignment operator, more explicit than `=`

Let's Start the Analysis!

Setup

- Install **R** and **RStudio**
- Download the **Repository**



OR

- Use this RStudio Cloud workspace: <https://rstudio.cloud/project/355872>

Package Dependencies

Install packages (if local setup)...

```
install.packages("sf")  
install.packages("tidycensus")  
install.packages("dplyr")  
install.packages("ggplot2")  
install.packages("tidyr")  
install.packages("purrr")  
install.packages("lwgeom")
```

... and load libraries

```
library(dplyr)  
library(ggplot2)  
library(tidyr)  
library(sf)  
library(tidycensus)  
library(tigris)  
library(purrr)
```

Your Choices!

Get an API Key from http://api.census.gov/data/key_signup.html

```
census_api_key("<YOUR API KEY>")
demo_variables <- # define the variables you want to analyze here
de_census_data <- get_acs(geography = "tract",
                          state = "DE",
                          variables = demo_variables,
                          geometry = TRUE,
                          cb = TRUE)
```

OR

Load de_census_data.RData

```
load("data/de_census_data.RData")
```

Look at the data

```
head(de_census_data)
```

```
## Simple feature collection with 6 features and 5 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:           xmin: -75.7601 ymin: 39.17347 xmax: -75.60411 ymax: 39.29937
## epsg (SRID):    4269
## proj4string:     +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
##
```

	GEOID	NAME	variable	estimate	moe
## 1	10001040100	Census Tract 401, Kent County, Delaware	white	6080	502
## 2	10001040100	Census Tract 401, Kent County, Delaware	black	501	100
## 3	10001040100	Census Tract 401, Kent County, Delaware	asian	58	60
## 4	10001040100	Census Tract 401, Kent County, Delaware	hispanic	265	198
## 5	10001040100	Census Tract 401, Kent County, Delaware	foreignborn	132	125
## 6	10001040100	Census Tract 401, Kent County, Delaware	high_school_diplomas	1808	273

```
##
```

	geometry
## 1	MULTIPOLYGON (((-75.7601 39...
## 2	MULTIPOLYGON (((-75.7601 39...
## 3	MULTIPOLYGON (((-75.7601 39...
## 4	MULTIPOLYGON (((-75.7601 39...
## 5	MULTIPOLYGON (((-75.7601 39...
## 6	MULTIPOLYGON (((-75.7601 39...

Some cleaning using %>%

```
de_census_data_clean <- de_census_data %>%  
  separate(col = NAME,  
           into = c("Census_Tract", "County", "State"),  
           sep = ",") %>%  
  separate(col = Census_Tract,  
           into = c(NA, NA, "Census_Tract_Number"),  
           sep = " ")
```


Using the grammar

Let's read the previous code step by step

```
de_census_data %>% # take the data, and then
  separate(col = NAME,
           into = c("Census_Tract", "County", "State"),
           sep = ",") # separate by ','
```

```
## Simple feature collection with 2616 features and 7 fields (with 12 geometries empty)
```

```
## geometry type:  MULTIPOLYGON
```

```
## dimension:      XY
```

```
## bbox:           xmin: -75.78866 ymin: 38.45101 xmax: -75.04894 ymax: 39.83901
```

```
## epsg (SRID):    4269
```

```
## proj4string:     +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
```

```
## First 10 features:
```

##	GEOID	Census_Tract	County	State	variable	estimate	moe
## 1	10001040100	Census Tract 401	Kent County	Delaware	white	6080	502
## 2	10001040100	Census Tract 401	Kent County	Delaware	black	501	100
## 3	10001040100	Census Tract 401	Kent County	Delaware	asian	58	60
## 4	10001040100	Census Tract 401	Kent County	Delaware	hispanic	265	198
## 5	10001040100	Census Tract 401	Kent County	Delaware	foreignborn	132	125
## 6	10001040100	Census Tract 401	Kent County	Delaware	high_school_diplomas	1808	273
## 7	10001040100	Census Tract 401	Kent County	Delaware	bachelor_degrees	268	122
## 8	10001040100	Census Tract 401	Kent County	Delaware	masters_degrees	181	97
## 9	10001040100	Census Tract 401	Kent County	Delaware	households_earning_over_200k	51	50
## 10	10001040100	Census Tract 401	Kent County	Delaware	median_income	63324	8985

```
## geometry
```

```
## 1 MULTIPOLYGON (((-75.7601 39...
```

```
## 2 MULTIPOLYGON (((-75.7601 39...
```

```
## 3 MULTIPOLYGON (((-75.7601 39...
```

```
## 4 MULTIPOLYGON (((-75.7601 39...
```

```
## 5 MULTIPOLYGON (((-75.7601 39...
```

```
## 6 MULTIPOLYGON (((-75.7601 39...
```

```
## 7 MULTIPOLYGON (((-75.7601 39...
```

```
## 8 MULTIPOLYGON (((-75.7601 39...
```

```
## 9 MULTIPOLYGON (((-75.7601 39...
```

```
## 10 MULTIPOLYGON (((-75.7601 39...
```

Using the grammar (continued)

```
de_census_data %>% # take the data, and then
  separate(col = NAME,
    into = c("Census_Tract", "County", "State"),
    sep = ",") %>% # separate, and then
  separate(col = Census_Tract,
    into = c(NA, NA, "Census_Tract_Number"),
    sep = " ") # separate out Number
```

```
## Simple feature collection with 2616 features and 7 fields (with 12 geometries empty)
```

```
## geometry type:  MULTIPOLYGON
```

```
## dimension:      XY
```

```
## bbox:           xmin: -75.78866 ymin: 38.45101 xmax: -75.04894 ymax: 39.83901
```

```
## epsg (SRID):    4269
```

```
## proj4string:     +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
```

```
## First 10 features:
```

##	GEOID	Census_Tract_Number	County	State	variable	estimate
## 1	10001040100	401	Kent County	Delaware	white	6080
## 2	10001040100	401	Kent County	Delaware	black	501
## 3	10001040100	401	Kent County	Delaware	asian	58
## 4	10001040100	401	Kent County	Delaware	hispanic	265
## 5	10001040100	401	Kent County	Delaware	foreignborn	132
## 6	10001040100	401	Kent County	Delaware	high_school_diplomas	1808
## 7	10001040100	401	Kent County	Delaware	bachelor_degrees	268
## 8	10001040100	401	Kent County	Delaware	masters_degrees	181
## 9	10001040100	401	Kent County	Delaware	households_earning_over_200k	51
## 10	10001040100	401	Kent County	Delaware	median_income	63324

```
## moe geometry
```

```
## 1 502 MULTIPOLYGON (((-75.7601 39...
```

```
## 2 100 MULTIPOLYGON (((-75.7601 39...
```

```
## 3 60 MULTIPOLYGON (((-75.7601 39...
```

```
## 4 198 MULTIPOLYGON (((-75.7601 39...
```

```
## 5 125 MULTIPOLYGON (((-75.7601 39...
```

```
## 6 273 MULTIPOLYGON (((-75.7601 39...
```

```
## 7 122 MULTIPOLYGON (((-75.7601 39...
```

```
## 8 97 MULTIPOLYGON (((-75.7601 39...
```

```
## 9 50 MULTIPOLYGON (((-75.7601 39...
```

```
## 10 8985 MULTIPOLYGON (((-75.7601 39...
```

Creating the Map!

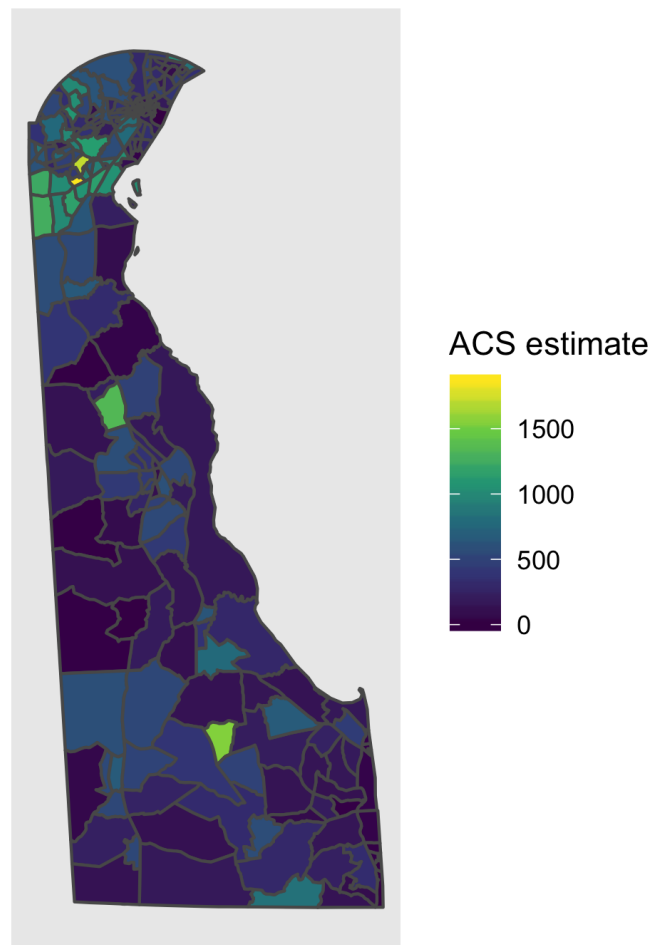
First pick one variable to view estimates

```
# Let's look at the number of foreignborn in each Tract  
## create a data frame by subsetting only the 'foreignborn'  
de_census_fb <- de_census_data_clean %>%  
  filter(variable %in% c("foreignborn"))
```

Now let's plot this:

```
ggplot(de_census_fb, aes(fill = estimate)) +  
  geom_sf() +  
  scale_fill_viridis_c() +  
  coord_sf(datum = NA) +  
  labs(title = "Foreign-Born Estimates by DE Census Tract",  
        caption = "Data: 2013-2017 5-year ACS",  
        fill = "ACS estimate")
```

Foreign-Born Estimates by DE Census Tract



Data: 2013-2017 5-year ACS


See the individual plot layers

We use the ggplot2 package for **layering** plot info. `geom_sf` is used to map the varied shapes (polygons, lines)

```
ggplot(de_census_fb, aes(fill = estimate))
```

Add the geometries

```
ggplot(de_census_fb, aes(fill = estimate)) +  
  geom_sf()
```



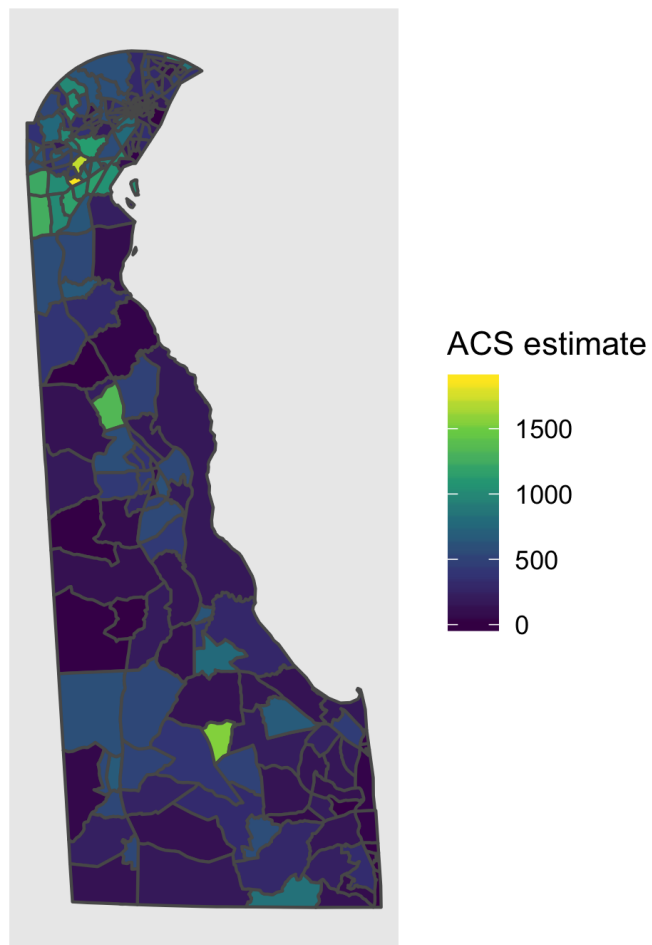
Add safe colors

```
ggplot(de_census_fb, aes(fill = estimate)) +  
  geom_sf() +  
  scale_fill_viridis_c()
```

Add the theme and title

```
ggplot(de_census_fb, aes(fill = estimate)) +  
  geom_sf() +  
  scale_fill_viridis_c() +  
  coord_sf(crs = 26916, datum = NA) +  
  labs(title = "Foreign-Born Estimates by DE Census Tract",  
        caption = "Data: 2013-2017 5-year ACS",  
        fill = "ACS estimate")
```


Foreign-Born Estimates by DE Census Tract



Data: 2013-2017 5-year ACS

Subset only Wilmington areas using Tract

```
# Create a dataframe with only Wilmington Tracts
wilm_census_data <- de_census_data_clean %>%
  filter(Census_Tract_Number %in% c(2, 3, 4, 5, 6.01, 6.02,
                                     9, 11, 12, 13, 14, 15,
                                     16, 19.02, 21, 22, 23, 24,
                                     25, 26, 27, 28, 29, 30.02))
```

Cleaning came of help here to filter out only the relevant numbers

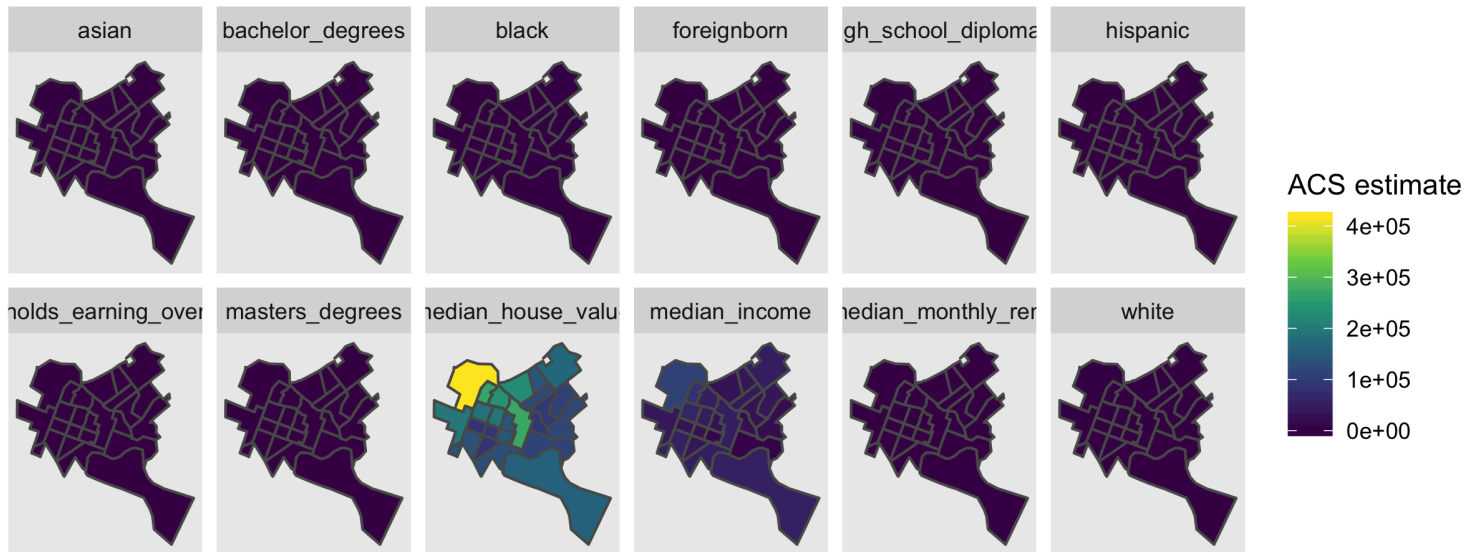
Let's plot EVERYTHING!

```
# Plot all our data
ggplot(wilm_census_data, aes(fill = estimate)) +
  geom_sf() +
  scale_fill_viridis_c() +
  coord_sf(crs = 26916, datum = NA) +
  labs(title = "Estimates by Census Tract",
        subtitle = "Wilmington, DE",
        caption = "Data: 2013-2017 5-year ACS
\nData acquired with the R tidycensus package.",
        fill = "ACS estimate") +
  facet_wrap(~variable)
```

facet helps you split up the data by variable and plot each

What's the problem?

Estimates by Census Tract
Wilmington, DE



Data: 2013-2017 5-year ACS

Data acquired with the R tidycensus package.

Plotting some comparable variables

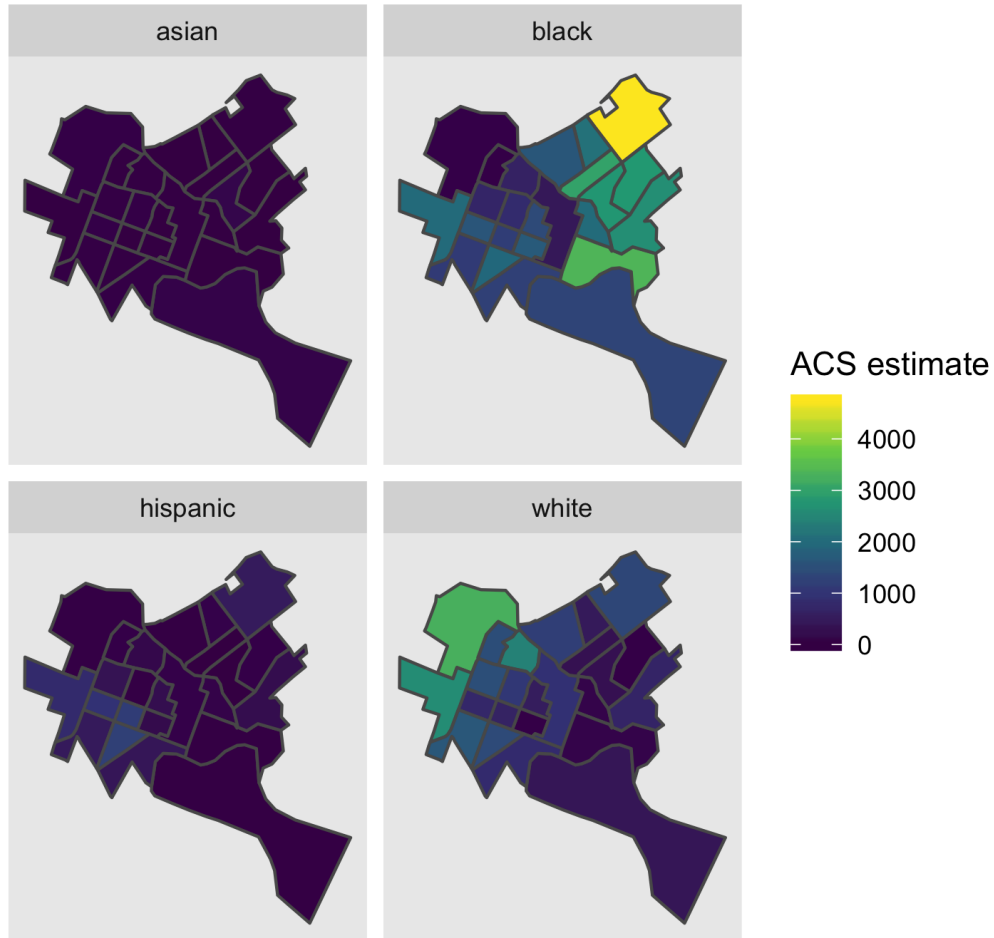
Let's focus on the race variables

```
## create a data frame with race variables  
wilm_census_race <- wilm_census_data %>%  
  filter(variable %in% c("hispanic", "black", "asian", "white"))
```

```
## plot  
ggplot(wilm_census_race, aes(fill = estimate)) +  
  geom_sf() +  
  scale_fill_viridis_c() +  
  coord_sf(crs = 26916, datum = NA) +  
  labs(title = "Population Estimates",  
        subtitle = "Wilmington, DE",  
        fill = "ACS estimate") +  
  facet_wrap(~variable)
```

Population Estimates

Wilmington, DE



Comparing a better way

'Small multiples plots' are useful to compare between variables. But we need to make sure we compare the right proportions so as to not let people take away a wrong insight.

Let's do some data aggregation and data joins to find the percentages within each tract.

Code

Create the total population data frame

```
## Estimate the total population  
## by summing up the different race estimates  
wilm_tract_pop <- wilm_census_race %>%  
  group_by(Census_Tract_Number) %>%  
  summarise(Population_Estimate = sum(estimate))
```

```
# remove geometry variable to make it a regular dataset  
st_geometry(wilm_tract_pop) <- NULL
```

Join the total population data with the original and create the Percentage column

```
# create data frame with the percentages  
wilm_tract_percpop <- wilm_census_race %>%  
  left_join(wilm_tract_pop, by = "Census_Tract_Number") %>%  
  mutate(Percentage = estimate/Population_Estimate)
```


Step by Step

Highlight sections and run using cmd + return to see the separate steps.
Remember not to highlight the variable assignment part.

```
wilm_census_race %>%
  group_by(Census_Tract_Number)

## Simple feature collection with 96 features and 7 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:           xmin: -75.5885 ymin: 39.7005 xmax: -75.51268 ymax: 39.77263
## epsg (SRID):    4269
## proj4string:     +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
## # A tibble: 96 x 8
## # Groups:   Census_Tract_Number [24]
##   GEOID Census_Tract_Num... County State variable estimate moe geometry
##   <chr> <chr>                <chr> <chr> <chr>      <dbl> <dbl> <MULTIPOLYGON [°]>
## 1 10003... 2                " New C... " De... white      1343  288 (((-75.53885 39.76623, -75.53547...
## 2 10003... 2                " New C... " De... black       4741  722 (((-75.53885 39.76623, -75.53547...
## 3 10003... 2                " New C... " De... asian         32   60 (((-75.53885 39.76623, -75.53547...
## 4 10003... 2                " New C... " De... hispanic      505  342 (((-75.53885 39.76623, -75.53547...
## 5 10003... 3                " New C... " De... white       468  133 (((-75.54275 39.76461, -75.53923...
## 6 10003... 3                " New C... " De... black      2111  244 (((-75.54275 39.76461, -75.53923...
## 7 10003... 3                " New C... " De... asian         0   11 (((-75.54275 39.76461, -75.53923...
## 8 10003... 3                " New C... " De... hispanic      147  121 (((-75.54275 39.76461, -75.53923...
## 9 10003... 4                " New C... " De... white      1197  209 (((-75.55714 39.76163, -75.55692...
## 10 10003... 4               " New C... " De... black      1645  291 (((-75.55714 39.76163, -75.55692...
## # ... with 86 more rows
```

```
wilm_census_race %>%
  group_by(Census_Tract_Number) %>%
  summarise(Population_Estimate = sum(estimate))
```

```
## Simple feature collection with 24 features and 2 fields
## geometry type: MULTIPOLYGON
## dimension: XY
## bbox: xmin: -75.5885 ymin: 39.7005 xmax: -75.51268 ymax: 39.77263
## epsg (SRID): 4269
## proj4string: +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
## # A tibble: 24 x 3
##   Census_Tract_Numb... Population_Estima... geometry
##   <chr> <dbl> <MULTIPOLYGON [°]>
## 1 11 3267 (((-75.56156 39.75696, -75.56049 39.75772, -75.56151 39.75...
## 2 12 1691 (((-75.56816 39.75632, -75.56531 39.76056, -75.56244 39.76...
## 3 13 3508 (((-75.57996 39.76337, -75.57785 39.76644, -75.5768 39.768...
## 4 14 2601 (((-75.57097 39.75227, -75.57013 39.75353, -75.56863 39.75...
## 5 15 2017 (((-75.56376 39.7493, -75.56289 39.75057, -75.56229 39.751...
## 6 16 2223 (((-75.55931 39.74559, -75.55808 39.747, -75.55719 39.7482...
## 7 19.02 1879 (((-75.56028 39.73048, -75.55854 39.73517, -75.55634 39.73...
## 8 2 6621 (((-75.53885 39.76623, -75.53547 39.76727, -75.53466 39.76...
## 9 21 2044 (((-75.56197 39.74169, -75.56072 39.74348, -75.56032 39.74...
## 10 22 2990 (((-75.5662 39.74573, -75.56479 39.74781, -75.55931 39.745...
## # ... with 14 more rows
```

```
wilm_census_race %>%
  left_join(wilm_tract_pop, by = "Census_Tract_Number")
```

```
## Simple feature collection with 96 features and 8 fields
## geometry type: MULTIPOLYGON
## dimension: XY
## bbox: xmin: -75.5885 ymin: 39.7005 xmax: -75.51268 ymax: 39.77263
## epsg (SRID): 4269
## proj4string: +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
## First 10 features:
```

##	GEOID	Census_Tract_Number	County	State	variable	estimate	moe
## 1	10003000200	2	New Castle County	Delaware	white	1343	288
## 2	10003000200	2	New Castle County	Delaware	black	4741	722
## 3	10003000200	2	New Castle County	Delaware	asian	32	60
## 4	10003000200	2	New Castle County	Delaware	hispanic	505	342
## 5	10003000300	3	New Castle County	Delaware	white	468	133
## 6	10003000300	3	New Castle County	Delaware	black	2111	244
## 7	10003000300	3	New Castle County	Delaware	asian	0	11
## 8	10003000300	3	New Castle County	Delaware	hispanic	147	121
## 9	10003000400	4	New Castle County	Delaware	white	1197	209
## 10	10003000400	4	New Castle County	Delaware	black	1645	291

```
## Population_Estimate geometry
## 1 6621 MULTIPOLYGON (((-75.53885 3...
## 2 6621 MULTIPOLYGON (((-75.53885 3...
## 3 6621 MULTIPOLYGON (((-75.53885 3...
## 4 6621 MULTIPOLYGON (((-75.53885 3...
## 5 2726 MULTIPOLYGON (((-75.54275 3...
## 6 2726 MULTIPOLYGON (((-75.54275 3...
## 7 2726 MULTIPOLYGON (((-75.54275 3...
## 8 2726 MULTIPOLYGON (((-75.54275 3...
## 9 2890 MULTIPOLYGON (((-75.55714 3...
## 10 2890 MULTIPOLYGON (((-75.55714 3...
```

```
wilm_census_race %>%
  left_join(wilm_tract_pop, by = "Census_Tract_Number") %>%
  mutate(Percentage = estimate/Population_Estimate)
```

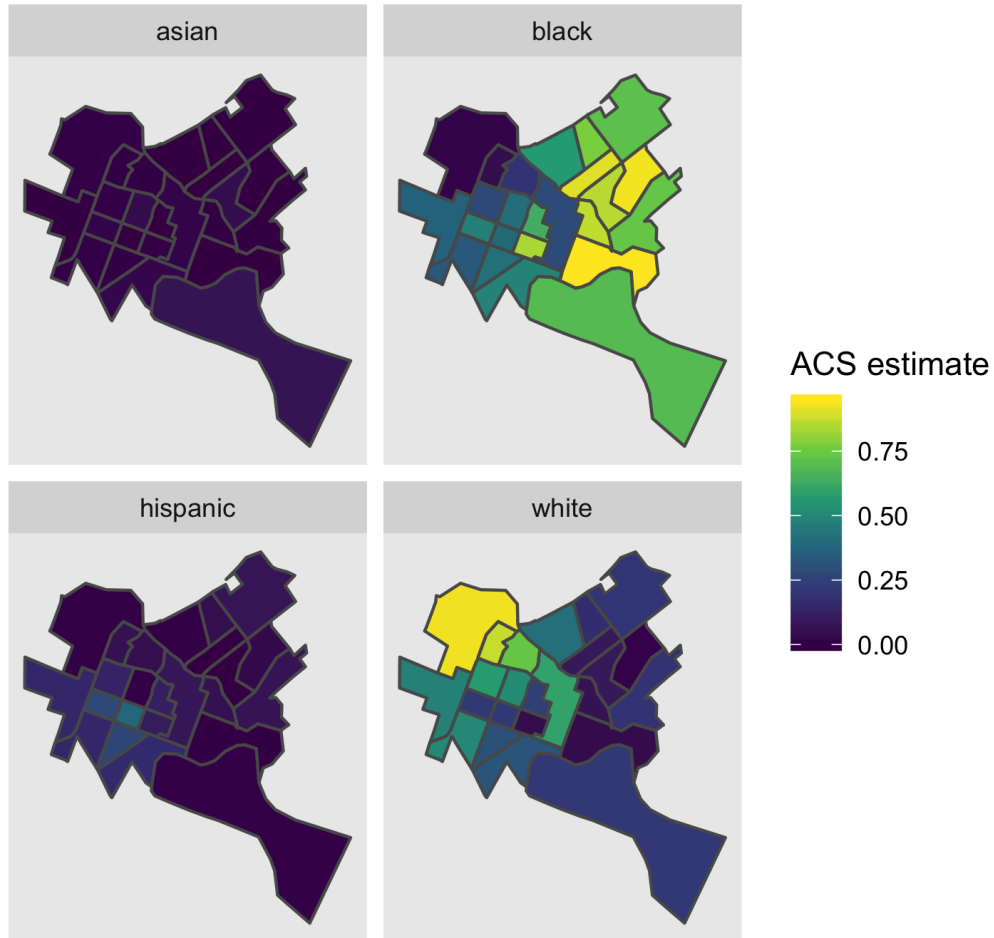
```
## Simple feature collection with 96 features and 9 fields
## geometry type: MULTIPOLYGON
## dimension: XY
## bbox: xmin: -75.5885 ymin: 39.7005 xmax: -75.51268 ymax: 39.77263
## epsg (SRID): 4269
## proj4string: +proj=longlat +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +no_defs
## First 10 features:
##      GEOID Census_Tract_Number County State variable estimate moe
## 1 10003000200 2 New Castle County Delaware white 1343 288
## 2 10003000200 2 New Castle County Delaware black 4741 722
## 3 10003000200 2 New Castle County Delaware asian 32 60
## 4 10003000200 2 New Castle County Delaware hispanic 505 342
## 5 10003000300 3 New Castle County Delaware white 468 133
## 6 10003000300 3 New Castle County Delaware black 2111 244
## 7 10003000300 3 New Castle County Delaware asian 0 11
## 8 10003000300 3 New Castle County Delaware hispanic 147 121
## 9 10003000400 4 New Castle County Delaware white 1197 209
## 10 10003000400 4 New Castle County Delaware black 1645 291
##      Population_Estimate geometry Percentage
## 1 6621 MULTIPOLYGON (((-75.53885 3... 0.202839450
## 2 6621 MULTIPOLYGON (((-75.53885 3... 0.716054977
## 3 6621 MULTIPOLYGON (((-75.53885 3... 0.004833107
## 4 6621 MULTIPOLYGON (((-75.53885 3... 0.076272466
## 5 2726 MULTIPOLYGON (((-75.54275 3... 0.171680117
## 6 2726 MULTIPOLYGON (((-75.54275 3... 0.774394718
## 7 2726 MULTIPOLYGON (((-75.54275 3... 0.000000000
## 8 2726 MULTIPOLYGON (((-75.54275 3... 0.053925165
## 9 2890 MULTIPOLYGON (((-75.55714 3... 0.414186851
## 10 2890 MULTIPOLYGON (((-75.55714 3... 0.569204152
```

Note: tidycensus allows you to get the summary value through the API as well!

Plot the percentage

```
ggplot(wilm_tract_percpop, aes(fill = Percentage)) +  
  geom_sf() +  
  scale_fill_viridis_c() +  
  coord_sf(crs = 26916, datum = NA) +  
  labs(title = "Percentage of Total (Estimates) by Census Tract",  
        subtitle = "Wilmington, DE",  
        fill = "ACS estimate") +  
  facet_wrap(~variable)
```

Percentage of Total (Estimates) by Census Tract Wilmington, DE



Your Turn

You can create a Wilmington Education variable by filtering `c("high_school_diplomas", "bachelor_degrees", "masters_degrees")` variables and recreating the previous visualizations:

- Count estimates by census tract
- Percentage estimates by census tract

Answer

```
wilm_census_edu <- wilm_census_data %>%  
  filter(variable %in% c("high_school_diplomas",  
                        "bachelor_degrees",  
                        "masters_degrees"))  
  
ggplot(wilm_census_edu, aes(fill = estimate)) +  
  geom_sf() +  
  scale_fill_viridis_c() +  
  coord_sf(crs = 26916, datum = NA) +  
  labs(title = "Education Estimates",  
       subtitle = "Wilmington, DE",  
       fill = "ACS estimate") +  
  facet_wrap(~variable)
```



```

# As a percentage of education data available
# (or use total population)
wilm_tract_totedu <- wilm_census_edu %>%
  group_by(Census_Tract_Number) %>%
  summarise(Education_Estimate = sum(estimate))

st_geometry(wilm_tract_totedu) <- NULL # remove geometry

wilm_tract_percedu <- wilm_tract_totedu %>%
  left_join(wilm_census_edu, by = "Census_Tract_Number") %>%
  mutate(Percentage = estimate/Education_Estimate)

# Plot
ggplot(wilm_tract_percedu, aes(fill = Percentage)) +
  geom_sf() +
  scale_fill_viridis_c() +
  coord_sf(crs = 26916, datum = NA) +
  labs(title = "Percentage of Education Estimates by Census Tract",
        subtitle = "Wilmington, DE",
        fill = "ACS estimate") +
  facet_wrap(~variable)

```

Dot Density Plots

Choropleth maps have a tendency of being misunderstood due to the *area* covered by a color. We can plot dots in order to avoid the issue of misrepresentation of sparsely populated areas and give an idea of density.

Hold tight as this will have some heavy lifting with functions from dplyr and purrr!

```
wm_dots <- map(c("white", "black",  
                "asian", "hispanic"), function(group) {  
  wilm_census_data %>%  
    filter(variable == group) %>%  
    st_sample(., size = .$estimate / 10, exact = FALSE) %>%  
    st_sf() %>%  
    mutate(group = group)  
}) %>%  
  reduce(rbind) %>%  
  group_by(group) %>%  
  summarize()
```


`map()` Applies a function to each element of a vector, in our case the vector is the race values. i.e, for each race we subset the wilmington data, and create dots representing the population in each tract.

```
# for each group
group <- "asian"
wilm_census_data %>%
  filter(variable == group) %>%
  st_sample(., size = .$estimate / 10, exact = FALSE) %>%
  st_sf() %>%
  mutate(group = group)
```

- `st_sample()` generates a sample of random dots each one representing 10 people.
- `st_sf()` converts the POINT geometry set back to simple features dataframe.

Plotting it all together

```
ggplot() +  
  geom_sf(data = wilm_census_data, color = "grey95", fill = "white")  
  geom_sf(data = wm_dots, aes(color = group, fill = group),  
    size = 0.1, alpha = 0.5) +  
  theme_minimal()
```



Looking back

Data being tidy allowed us to immediately use commands like:

- `facet`
- `group_by`

Remember ACS are estimates so we should consider the MOE or Margin of Error variable.

Next steps

- Try using another layer with alpha to show the MOE *
- Integrate this data with statistics from the Uniform Crime Reporting database
- Ask interesting questions and use tidy functions to get quick results
 - What are the high median rent areas?
 - Does the median house value correlate with houses earning above 200k?
 - Among the high median house value ones what is the percentage of owners to renter?
 - Does more Education level imply more Median income?
 - Where are the Vacant houses and what is the median income in these areas?

Thank You!